

M7 Characteristics of Time Series

General Insurance Modelling : Actuarial Modelling III ¹

Professor Benjamin Avanzi



24 April 2023

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)

1 Introduction (TS 1.0)

- Definition
- Applications
- Process for time series analysis

Definition

- Consider data of the same nature that have been observed at different points in time
- The mere fact that they are of the same nature means that they are likely related in one way or another - let's call those 'correlations' (an acceptable term in this context as we focus on this measure, at least in this course)
- This is in contrast with the usual "i.i.d." assumptions associated with a sample of outcomes of a random variable
- This invalidates some of the techniques we know, and brings additional difficulties, but also opportunities! (such as forecasting)

Definition: "The systematic approach by which one goes about answering the mathematical and statistical questions posed by these time correlations is commonly referred to as **time series analysis**."

1 Introduction (TS 1.0)

- Definition
- Applications
- Process for time series analysis

Applications

The applications of time series are many, and crucial in many cases:

- Economics: unemployment, GDP, CPI, etc ...
- Finance: share prices, indices, etc ...
- Medicine: COVID-19 cases and fatalities, biometric data for a patient (blood pressure, iron levels, ...), etc ...
- Global warming: ocean temperatures, CO₂ levels, particule levels in the atmosphere, sea levels, all in relation with another, and with many others
- **Actuarial studies:** frequency and severity of claims in a LoB, mortality (at different ages, in different locations, ...), superimposed inflation, IBNR claims, etc ...

1 Introduction (TS 1.0)

- Definition
- Applications
- Process for time series analysis

Process for time series analysis

Sketch of process:

- Careful examination of data plotted over time (Module 7)
- Compute major statistical indicators (Modules 7 and 8)
- Guess an appropriate method for analysing the data (Modules 8 and 9)
- Fit and assess your model (Module 9)
- Use your model to perform forecasts if relevant (Module 10)

We distinguish two types of approaches:

- *Time domain approach*: investigate lagged relationships (impact of today on tomorrow)
- *Frequency domain approach*: investigate cycles (understand regular variations)

In actuarial studies, both are relevant.


- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)**
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)

2 Examples (TS 1.1)

- Johnson & Johnson quarterly earnings per share
- Global mean land-ocean temperature index
- Dow Jones Industrial Average
- Analysis of two series together: El Niño & fish population
- Signals within noise

Johnson & Johnson quarterly earnings per share

Johnson & Johnson quarterly earnings per share,
84 quarters, 1960–I to 1980–IV



What is the primary pattern?


Can you see any cyclical pattern as well?

How does volatility change over time (if at all)?

2 Examples (TS 1.1)

- Johnson & Johnson quarterly earnings per share
- Global mean land-ocean temperature index
- Dow Jones Industrial Average
- Analysis of two series together: El Niño & fish population
- Signals within noise

Global mean land-ocean temperature index




Can you see a trend? Are there periods of continuous increase?
What would be the main focus for global warming: trend or cycles?
How does this graph support the global warming thesis?

2 Examples (TS 1.1)

- Johnson & Johnson quarterly earnings per share
- Global mean land-ocean temperature index
- **Dow Jones Industrial Average**
- Analysis of two series together: El Niño & fish population
- Signals within noise

Dow Jones Industrial Average




How is this time series special?

What qualities would a good forecast model need to have?

2 Examples (TS 1.1)

- Johnson & Johnson quarterly earnings per share
- Global mean land-ocean temperature index
- Dow Jones Industrial Average
- Analysis of two series together: El Niño & fish population
- Signals within noise

Analysis of two series together: El Niño & fish population




How many cycles can you spot?

Is there a relationship between both series?

2 Examples (TS 1.1)

- Johnson & Johnson quarterly earnings per share
- Global mean land-ocean temperature index
- Dow Jones Industrial Average
- Analysis of two series together: El Niño & fish population
- Signals within noise

Signals within noise



Typically we only see the the signal obscured by noise.

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)**
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)

3 Basic models (TS 1.2)

- Preliminaries
 - White noise - 3 scales :-)
 - Gaussian white noise series and its 3-point moving average
 - Filtering (and moving average)
 - Autoregressions
 - Autoregression examples
 - Random walk with drift

Preliminaries

- Our primary objective is to develop mathematical models that provide plausible descriptions for sample data.
- A time series is a sequence of rv's x_1, x_2, x_3, \dots , **denoted** $\{x_t\}$
- In this course, t will typically be discrete and be $\in \mathbb{N}$ (or subset)
- One set of observed values of $\{x_t\}$ is referred to as a *realisation*
- Time series are usually plotted with time in the x -axis, with observations connected at adjacent periods
- Sampling rate must be sufficient, lest appearance of the data is changed completely (*aliasing*; see also [this](#) which explains how car wheels can appear to go backwards)
- Smoothness of the time series suggests some level of correlation between adjacent points, or in other words that x_t depends in some way on the past values $x_{t-1}, x_{t-2}, \dots \rightarrow$ This is a good starting point for imagining appropriate theoretical models!

3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- Filtering (and moving average)
- Autoregressions
- Autoregression examples
- Random walk with drift

White noise - 3 scales :-)

Let's define w_t as **uncorrelated** random variables w_t with mean 0 and finite variance σ_w^2 . This is denoted

$$w_t \sim \text{wn}(0, \sigma_w^2),$$

and is called a **white noise**. Two special cases:

- **White independent noise:** (or iid noise) additional assumption of *iid*, denoted

$$w_t \sim \text{iid}(0, \sigma_w^2).$$

- **Gaussian white noise:** further additional assumption of *normal distribution*, denoted

$$w_t \sim \text{iid } N(0, \sigma_w^2).$$


Usually, time series are smoother than that (see bottom graph on the next slide). Ways of introducing *serial correlation* and more *smoothness* into time series include **filtering** and **autoregression**.

3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- Filtering (and moving average)
- Autoregressions
- Autoregression examples
- Random walk with drift

Gaussian white noise series and its 3-point moving average

```
w <- rnorm(500, 0, 1) # 500  $N(0,1)$  variates  
plot.ts(w, ylim = c(-3, 3), main = "white noise")  
v <- stats::filter(w, sides = 2, filter = rep(1/3, 3)) # moving average  
plot.ts(v, ylim = c(-3, 3), main = "moving average")
```



3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- **Filtering (and moving average)**
- Autoregressions
- Autoregression examples
- Random walk with drift

Filtering (and moving average)

A series v_t which is a **linear combination** of values of a more fundamental time series w_t is called a **filtered series**.

- Example: 3-point moving average (see bottom of previous slide for graph):

$$v_t = \frac{1}{3}(w_{t-1} + w_t + w_{t+1}).$$

- In R, moving averages are implemented through the function

```
filter(x, filter, method = c("convolution"), sides = 2)
```

where `x` is the original series, `filter` is a vector of weights (in reverse time order), `method = c("convolution")` is the default (alternative is `recursive`), and where `sides` is 1 for past values only, and 2 if weights are centered around lag 0 (requires uneven number of weights).

- Moving average smoothers will be further discussed in Module 8.

3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- Filtering (and moving average)
- **Autoregressions**
- Autoregression examples
- Random walk with drift

Autoregressions

A series x_t that depends on some of its past values, as well as a noise w_t is called an **autoregression**, because the formula looks like a regression—not of independent variables, but of its own past values—hence **autoregression**.

- Example: An autoregression of the white noise:

$$x_t = x_{t-1} - 0.9x_{t-2} + w_t.$$

- If the autoregression goes back k periods, one needs k initial conditions (filter will use 0's otherwise).
- In R, autoregressions are implemented through the function

```
filter(x, filter, method = c("recursive"), init)
```

where x is the original series, $filter$ is a vector of weights (reverse time order) and $init$ a vector of initial values (reverse time order).



- Autoregressions will be denoted $AR(p)$ (details in Module 9).

3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- Filtering (and moving average)
- Autoregressions
- **Autoregression examples**
- Random walk with drift

Autoregression examples

```
w = rnorm(550,0,1) # 50 extra to avoid startup problems
x = stats::filter(w, filter=c(1,-.9), method="recursive")[-(1:50)]
# remove first 50
plot.ts(x, ylab="autogression", #
        main="Autoregressive series generated from model  $x_t = x_{t-1} - 0.9x_{t-1}$ ")
y = stats::filter(w, filter=c(1,-.3), method="recursive")[-(1:50)]
# remove first 50
plot.ts(y, ylab="autogression", #
        main="Autoregressive series generated from model  $x_t = x_{t-1} - 0.3x_{t-1}$ ")
```

Autoregressive series generated from model $x_t = x_{t-1} - 0.9x_{t-2} + w_t$ Autoregressive series generated from model $x_t = x_{t-1} - 0.3x_{t-2} + w_t$ 

3 Basic models (TS 1.2)

- Preliminaries
- White noise - 3 scales :-)
- Gaussian white noise series and its 3-point moving average
- Filtering (and moving average)
- Autoregressions
- Autoregression examples
- Random walk with drift

Random walk with drift


- The autoregressions introduced above are all centered around 0 for all t (in the expected sense).
- Assume now that the series increases linearly by δ (called **drift**) every time unit.
- The **random walk with drift** looks back only one time unit:

$$x_t = \delta + x_{t-1} + w_t = \delta t + \sum_{j=1}^t w_j \text{ for } t = 1, 2, \dots$$

with initial condition $x_0 = 0$ and with w_t a white noise.

- If $\delta = 0$ this is simply called a **random walk**.
- The term can be explained by visualising each increment from t to $t + 1$ as a purely random step from wherever the process is at x_t , ignoring what happened before.

Random walk with drift $\delta = 0.2$ and $\sigma_w = 1$



Code used to generate the plot:

```
set.seed(155)  # so you can reproduce the results
w <- rnorm(200)
x <- cumsum(w)
wd <- w + 0.2
xd <- cumsum(wd)
plot.ts(xd, ylim = c(-5, 55), main = "random walk", ylab = "")
lines(x, col = 4)
abline(h = 0, col = 4, lty = 2)
```

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)**
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)

4 Describing the behaviour of basic models (TS 1.3)

- Motivation
- Full specification
- Mean function
- Autocovariance function
- The autocorrelation function (ACF)

Motivation

- In this section we would like to develop **theoretical measures** to help describe how times series behave.
- We are particularly interested in describing the **relationships** between observations at different points in time.

4 Describing the behaviour of basic models (TS 1.3)

- Motivation
- Full specification
- Mean function
- Autocovariance function
- The autocorrelation function (ACF)

Full specification

- A full specification of a time series of size n at times t_1, t_2, \dots, t_n for any n would require the full joint distribution function

$$F_{t_1, t_2, \dots, t_n}(c_1, c_2, \dots, c_n) = \Pr[x_{t_1} \leq c_1, x_{t_2} \leq c_2, \dots, x_{t_n} \leq c_n].$$

This is a quite unwieldy tool for analysis.

- Examination of the margins $F_t(x) = \Pr[x_t \leq x]$ and corresponding pdf $f_t(x)$, *when they exist*, can be informative.
- These are very theoretical. In practice, one often have only **one** realisation for each x_t so that inferring full distributions (let alone their dependence structure) is simply impractical without tricks, manipulations, and assumptions (some of which we will learn).

4 Describing the behaviour of basic models (TS 1.3)

- Motivation
- Full specification
- **Mean function**
- Autocovariance function
- The autocorrelation function (ACF)

Mean function

The **mean function** is defined as

$$\mu_{xt} = E[x_t] = \int_{-\infty}^{\infty} x f_t(x) dx.$$

Examples:

- **Moving Average Series:** we have

$$\mu_{vt} = E[v_t] = \frac{1}{3} (E[w_{t-1}] + E[w_t] + E[w_{t+1}]) = 0.$$

Smoothing does not change the mean.

- **Random walk with drift:** we have

$$\mu_{xt} = E[x_t] = \delta t + \sum_{j=1}^t E[w_j] = \delta t.$$

4 Describing the behaviour of basic models (TS 1.3)

- Motivation
- Full specification
- Mean function
- **Autocovariance function**
- The autocorrelation function (ACF)

Autocovariance function

The **autocovariance function** is defined as the second moment product

$$\gamma_x(s, t) = \text{Cov}(x_s, x_t) = E[(x_s - \mu_{xs})(x_t - \mu_{xt})]$$

for all s and t . Note:

- We will write $\gamma_x(s, t) = \gamma(s, t)$ if no confusion is possible.
- This is a measure of **linear** dependence.
- Smooth series \rightarrow large γ even for t and s far apart
- Choppy series $\rightarrow \gamma$ is nearly zero for large separations
- $[\gamma_x(s, t) = 0 \implies \text{independence}] \iff \text{all variables are normal}$

For two series x_t and y_t this becomes

$$\gamma_{xy}(s, t) = \text{Cov}(x_s, y_t) = E[(x_s - \mu_{xs})(y_t - \mu_{yt})],$$

called **cross-covariance function**.

Examples of autocovariance functions

White noise: The white noise series w_t has $E[w_t] = 0$ and

$$\gamma_w(s, t) = \text{Cov}(w_s, w_t) = \begin{cases} \sigma_w^2 & s = t \\ 0 & s \neq t \end{cases}$$

Remember that if

$$U = \sum_{j=1}^m a_j X_j$$

and

$$V = \sum_{k=1}^r b_k Y_k$$

then

$$\text{Cov}(U, V) = \sum_{j=1}^m \sum_{k=1}^r a_j b_k \text{Cov}(X_j, Y_k).$$

This will be useful for computing γ of filtered series.

Moving average: A 3-point moving average v_t to the white noise series w_t has

$$\gamma_v(s, t) = \text{Cov}(v_s, v_t) = \begin{cases} \frac{3}{9}\sigma_w^2 & s = t \\ \frac{2}{9}\sigma_w^2 & |s - t| = 1 \\ \frac{1}{9}\sigma_w^2 & |s - t| = 2 \\ 0 & |s - t| > 2 \end{cases}$$

This only depends on the time separation lag **only**, and not on the absolute location along the series.

This is related to the concept of *weak stationarity* which will introduce later.

Random walk: For the random walk $x_t = \sum_{j=1}^t w_j$ we have

$$\begin{aligned}\gamma_x(s, t) &= \text{Cov}(x_s, x_t) \\ &= \text{Cov}\left(\sum_{j=1}^s w_j, \sum_{k=1}^t w_k\right) \\ &= \min\{s, t\}\sigma_w^2.\end{aligned}$$

Contrary to the previous examples, this depends on the absolute location rather than the lag.

Also $\text{Var}(x_t) = t\sigma_w^2$ increases without bound as t increases.

4 Describing the behaviour of basic models (TS 1.3)

- Motivation
- Full specification
- Mean function
- Autocovariance function
- The autocorrelation function (ACF)

The autocorrelation function (ACF)

The **autocorrelation function (ACF)** is defined as

$$-1 \leq \rho(s, t) = \frac{\gamma(s, t)}{\sqrt{\gamma(s, s)\gamma(t, t)}} \leq 1.$$

- The ACF measures the *linear* predictability of the series at time t , say x_t , using *only* the value x_s .
- If we could do that perfectly, then $\rho(s, t) \pm 1$ and

$$x_t = \beta_0 + \beta_1 x_s$$

with β_1 of same sign as $\rho(s, t)$.

In the case of two series this becomes

$$-1 \leq \rho_{xy}(s, t) = \frac{\gamma_{xy}(s, t)}{\sqrt{\gamma_x(s, s)\gamma_y(t, t)}} \leq 1,$$

called **cross-correlation function (CCF)**.

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)**
- 6 Estimation of correlation (TS 1.5)

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- Examples of (non)-stationarity
- Trend stationarity
- Joint stationarity
- Example of joint stationarity
- Linear process

Strict stationarity

A **strictly stationary** times series is one for which the probabilistic behaviour of every collection of values $\{x_{t_1}, x_{t_2}, \dots, x_{t_k}\}$ is identical to that of the time shifted set (for any h) $\{x_{t_1+h}, x_{t_2+h}, \dots, x_{t_k+h}\}$. That is,

$$\begin{aligned}\Pr[x_{t_1} \leq c_1, x_{t_2} \leq c_2, \dots, x_{t_k} \leq c_k] \\ = \Pr[x_{t_1+h} \leq c_1, x_{t_2+h} \leq c_2, \dots, x_{t_k+h} \leq c_k]\end{aligned}$$

for all $k = 1, 2, \dots$, all the time points t_1, t_2, \dots, t_k , all numbers c_1, c_2, \dots, c_k and all time shifts $h = 0, \pm 1, \pm 2, \dots$. This implies

- identical marginals of dimensions $< k$ for any shift h
- constant mean: $\mu_{xs} = \mu_{xt} \equiv \mu$
- for $k = 2$, an autocovariance function that depends only on $t - s$:
 $\gamma(s, t) = \gamma(s + h, t + h)$

We need something less constraining, that still allows us to infer properties from a single series.

5 Stationary time series (TS 1.4)

- Strict stationarity
- **Weak stationarity**
- Properties of stationary series
- Examples of (non)-stationarity
- Trend stationarity
- Joint stationarity
- Example of joint stationarity
- Linear process

Weak stationarity

A **weakly stationary time series**, x_t , is a finite variance process such that

- 1 the mean value function, μ_{x_t} is **constant** and does not depend on time t , and
- 2 the autocovariance function, $\gamma(s, t)$ depends on s and t **only through their difference** $|s - t|$.

Note:

- We dropped full distributional requirements. This imposes conditions on the first two moments of the series only.
- Since those completely define a normal distribution, a (weak) stationary Gaussian time series is also strictly stationary.
- We will use the term stationary to mean weakly stationary; if a process is stationary in the strict sense, we will use the term strictly stationary.

Stationarity means we can estimate those two quantities by averaging of a *single* series. This is what we needed.

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- **Properties of stationary series**
- Examples of (non)-stationarity
- Trend stationarity
- Joint stationarity
- Example of joint stationarity
- Linear process

Properties of stationary series

- Because of condition 1,

$$\mu_t = \mu.$$

- Because of condition 2,

$$\gamma(t+h, t) = \text{Cov}(x_{t+h}, x_t) = \text{Cov}(x_h, x_0) = \gamma(h, 0) \equiv \gamma(h)$$

and the autocovariance of a stationary time series is then

$$\gamma(h) = \text{Cov}(x_{t+h}, x_t) = E[(x_{t+h} - \mu)(x_t - \mu)].$$

- $\gamma(h)$ is *non-negative definite*, which means that the variance of linear combinations of variates x_t will never be negative, that is,

$$0 \leq \text{Var}(a_1x_1 + \cdots + a_nx_n) = \sum_{j=1}^n \sum_{k=1}^n a_j a_k \gamma(j-k).$$

- Furthermore,

$$|\gamma(h)| \leq \gamma(0) \text{ (the variance of the time series)}$$

and

$$\gamma(h) = \gamma(-h).$$

- The autocorrelation function (ACF) of a stationary time series becomes

$$-1 \leq \rho(h) = \frac{\gamma(t+h, t)}{\sqrt{\gamma(t+h, t+h)\gamma(t, t)}} = \frac{\gamma(h)}{\gamma(0)} \leq 1.$$

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- **Examples of (non)-stationarity**
- Trend stationarity
- Joint stationarity
- Example of joint stationarity
- Linear process

Examples of (non)-stationarity

White noise: We have

$$\mu_{wt} = 0$$

and

$$\gamma_w(h) = \text{Cov}(w_{t+h}, w_t) = \begin{cases} \sigma_w^2 & h = 0, \\ 0 & h \neq 0, \end{cases},$$

which are both independent of time. Hence, the white noise satisfies both conditions and is (weakly) stationary. Furthermore,

$$\rho_w(h) = \begin{cases} 1 & h = 0, \\ 0 & h \neq 0. \end{cases}$$

If in addition $w_t \sim \text{iid } N(0, \sigma_w^2)$, then it is also strictly stationary.

Moving average: For the 3-point MA we have

$$\mu_{vt} = 0 \quad \text{and} \quad \gamma_v(h) = \begin{cases} \frac{3}{9}\sigma_w^2 & h = 0, \\ \frac{2}{9}\sigma_w^2 & h \pm 1, \\ \frac{1}{9}\sigma_w^2 & h \pm 2, \\ 0 & |h| > 2, \end{cases}$$

which are both independent of time. Hence, the 3-point MA satisfies both conditions and is stationary. Furthermore,

$$\rho_v(h) = \begin{cases} 1 & h = 0, \\ \frac{2}{3} & h \pm 1, \\ \frac{1}{3} & h \pm 2, \\ 0 & |h| > 2, \end{cases}$$

which is symmetric around lag 0.

Random walk: For the random walk model $x_t = \sum_{j=1}^t w_j$ we have

$$\mu_{xt} = \delta t,$$

which is a function of time t , and

$$\gamma(s, t) = \min\{s, t\}\sigma_w^2,$$

which depends on s and t (not just their difference), so the random walk is **not** stationary.

Furthermore, remember

$$\text{Var}(x_t) = \gamma_x(t, t) = t\sigma_w^2$$

which increases without bound as $t \rightarrow \infty$.

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- Examples of (non)-stationarity
- **Trend stationarity**
- Joint stationarity
- Example of joint stationarity
- Linear process

Trend stationarity

- If only the second condition (on the ACF) is satisfied, but not the first condition (on the mean value function), we have **trend stationarity**
- This means that the model has a stationary *behaviour* around its trend.
- Example: if

$$x_t = \alpha + \beta t + y_t \quad \text{where } y_t \text{ is stationary,}$$

then the mean function is

$$\mu_{x,t} = E[x_t] = \alpha + \beta t + \mu_y,$$

which is *not* independent of time. The autocovariance function,

$$\begin{aligned}\gamma_x(h) = \text{Cov}(x_{t+h}, x_t) &= E[(x_{t+h} - \mu_{x,t+h})(x_t - \mu_{x,t})] \\ &= E[(y_{t+h} - \mu_y)(y_t - \mu_y)] = \gamma_y(h)\end{aligned}$$

however, is independent of time.

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- Examples of (non)-stationarity
- Trend stationarity
- **Joint stationarity**
- Example of joint stationarity
- Linear process

Joint stationarity

Two time series, say, x_t and y_t , are said to be **jointly stationary** if they are each stationary, and the cross-covariance function

$$\gamma_{xy}(h) = \text{Cov}(x_{t+h}, y_t) = E[(x_{t+h} - \mu_x)(y_t - \mu_y)]$$

is a function only of lag h . The corresponding **cross-correlation function (CCF)** is

$$1 \leq \rho_{xy}(h) = \frac{\gamma_{xy}(h)}{\sqrt{\gamma_x(0)\gamma_y(0)}} \leq 1.$$

Note that because $\text{Cov}(x_2, y_1)$ and $\text{Cov}(x_1, y_2)$ (for example) need not be the same, it follows that typically

$$\rho_{xy}(h) \neq \rho_{xy}(-h),$$

that is, the CCF is **not generally symmetric about zero**. However, we have

$$\rho_{xy}(h) = \rho_{yx}(-h).$$

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- Examples of (non)-stationarity
- Trend stationarity
- Joint stationarity
- **Example of joint stationarity**
- Linear process

Example of joint stationarity

Consider

$$x_t = w_t + w_{t-1} \quad \text{and} \quad y_t = w_t - w_{t-1},$$

where w_t are independent with mean 0 and variance σ_w^2 . We have then

$$\begin{aligned} \gamma_x(0) &= \gamma_y(0) = 2\sigma_w^2 \\ \gamma_x(-1) &= \gamma_x(1) = \sigma_w^2 \\ \gamma_y(-1) &= \gamma_y(1) = -\sigma_w^2 \end{aligned}$$

and

$$\gamma_{xy}(-1) = -\sigma_w^2, \quad \gamma_{xy}(0) = 0, \quad \text{and} \quad \gamma_{xy}(1) = \sigma_w^2,$$

so that

$$\rho_{xy}(h) = \begin{cases} 0 & h = 0, \\ 1/2 & h = 1, \\ -1/2 & h = -1, \\ 0 & |h| \geq 2, \end{cases}$$

which depends only on the lag h , so both series are jointly stationary.

Prediction using cross-correlation

Prediction using cross-correlation: A lagging relation between two series x_t and y_t may be exploited for predictions. For instance, if

$$y_t = Ax_{t-\ell} + w_t,$$

x_t is said to *lead* y_t for $\ell > 0$, and is said to *lag* y_t for $\ell < 0$.

If the relation above holds true, then the lag ℓ can be inferred from the shape of the autocovariance of the input series x_t :

- If w_t is uncorrelated with x_t then

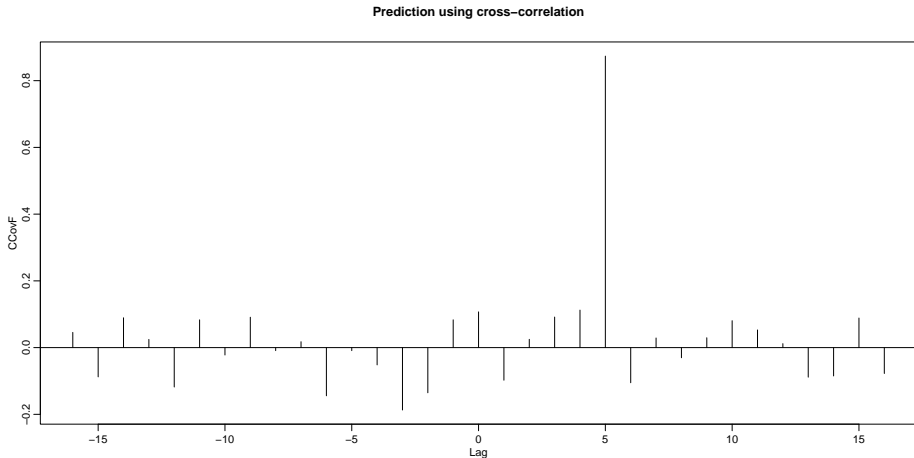
$$\begin{aligned}\gamma_{yx}(h) &= \text{Cov}(y_{t+h}, x_t) = \text{Cov}(Ax_{t+h-\ell} + w_{t+h}, x_t) \\ &= \text{Cov}(Ax_{t+h-\ell}, x_t) = A\gamma_x(h - \ell)\end{aligned}$$

- Since

$$\gamma_x(h - \ell) \leq \gamma_x(0),$$

the peak of $\gamma_{yx}(h)$ should be at $h = \ell$, and h will be positive if x_t leads y_t , negative if x_t lags y_t .

- Here $\ell = 5$ and x_t leads y_t :



Note this example was simulated and uses the R functions `lag` and `ccf`:

```
x <- rnorm(100)
y <- stats::lag(x, -5) + rnorm(100)
ccf(y, x, ylab = "CCovF", type = "covariance")
```

5 Stationary time series (TS 1.4)

- Strict stationarity
- Weak stationarity
- Properties of stationary series
- Examples of (non)-stationarity
- Trend stationarity
- Joint stationarity
- Example of joint stationarity
- Linear process

Linear process

A **linear process**, x_t , is defined to be a linear combination of white noise variates w_t , and is given by

$$x_t = \mu + \sum_{j=-\infty}^{\infty} \Psi_j w_{t-j}, \quad \sum_{j=-\infty}^{\infty} |\Psi_j| < \infty$$

This is an important class of models because it encompasses moving averages, autoregressions, and also the combination of both, called *autoregressive moving average (ARMA) processes* which we will introduce later.

Example:

- **Moving average** The 3-point moving average has

$$\Psi_0 = \Psi_{-1} = \Psi_1 = 1/3$$

and is hence a linear process.

Properties of linear processes

- The autocovariance function of a linear process is given by

$$\gamma_x(h) = \sigma_w^2 \sum_{j=-\infty}^{\infty} \Psi_{j+h} \Psi_j \quad \text{for } h \geq 0.$$

- It has finite variance if $\sum_{j=-\infty}^{\infty} \Psi_j^2 < \infty$.
- In its most general form x_t depends on the future ($j < 0$ components), the present ($j = 0$) and the past ($j > 0$).

For forecasting, a model dependent on the future is useless. We will focus on processes that do not depend on the future. Such processes are called **causal**, that is,

$$x_t \text{ is causal} \iff \Psi_j = 0 \quad \text{for } j < 0,$$

which we will assume unless stated otherwise.

- 1 Introduction (TS 1.0)
- 2 Examples (TS 1.1)
- 3 Basic models (TS 1.2)
- 4 Describing the behaviour of basic models (TS 1.3)
- 5 Stationary time series (TS 1.4)
- 6 Estimation of correlation (TS 1.5)**

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- Sample autocovariance function
- Sample autocorrelation function
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness

Background

- One can very rarely hypothesise (specify) time series. In practice, most analyses are performed using sample data.
- Furthermore, **one often has only one realisation of the time series**.
- This means that we don't have n realisations of the time series to estimate its covariance and correlation functions.
- This is why the **assumption of stationarity is essential**: in this case, the assumed 'homogeneity' of the data means we can estimate those functions on one realisation only.
- This also means that one needs to manipulate / de-trend series such that they are arguably stationary before we can fit parameters to them and use them for projections.

6 Estimation of correlation (TS 1.5)

- Background
- **Sample mean**
- Sample autocovariance function
- Sample autocorrelation function
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness

Sample mean

If a time series is stationary the mean function $\mu_t = \mu$ is constant so that we can estimate it by the sample mean,

$$\bar{x} = \frac{1}{n} \sum_{t=1}^n x_t.$$

This estimator is unbiased,

$$E[\bar{x}] = \mu,$$

and has standard error the square root of

$$\text{Var}(\bar{x}) = \frac{1}{n^2} \text{Cov} \left(\sum_{t=1}^n x_t, \sum_{s=1}^n x_s \right) = \frac{1}{n} \sum_{h=-n}^n \left(1 - \frac{|h|}{n} \right) \gamma_x(h).$$

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- **Sample autocovariance function**
- Sample autocorrelation function
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness

Sample autocovariance function

The **sample autocovariance function** is defined as

$$\hat{\gamma}(h) = \frac{1}{n} \sum_{t=1}^{n-h} (x_{t+h} - \bar{x})(x_t - \bar{x}) \quad \text{with } \hat{\gamma}(-h) = \hat{\gamma}(h) \text{ for } h = 0, 1, \dots, n-1.$$

Note:

- The estimator is biased.
- The sum runs over a restricted range $(n - h)$ because x_{t+h} is not available for $t + h > n$.
- One could wonder why the factor of the sum is not $1/(n - h)$ (the number of elements in the sum), but factor $1/n$ is not a mistake. It ensures that the estimate of the variances of linear combinations,

$$\widehat{\text{Var}}(a_1 x_1 + \dots + a_n x_n) = \sum_{j=1}^n \sum_{k=1}^n a_j a_k \hat{\gamma}(j - k),$$

is non-negative.

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- Sample autocovariance function
- **Sample autocorrelation function**
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness

Sample autocorrelation function

The **sample autocorrelation function (SACF)** is

$$\hat{\rho}(h) = \frac{\hat{\gamma}(h)}{\hat{\gamma}(0)}.$$

Under some conditions (see book for details), if x_t is a white noise, then for large n , the SACF $\hat{\rho}(h)$ is approximately normally distributed with zero mean and standard deviation given by

$$\sigma_{\hat{\rho}(h)} = \frac{1}{\sqrt{n}}.$$

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- Sample autocovariance function
- Sample autocorrelation function
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness


Testing for significance of autocorrelation

The asymptotic result for the variance of the SACF means we can test whether lagged observations are uncorrelated (which is a requirement for white noise):

- test for significance of the $\hat{\rho}$'s at different lags: check how many $\hat{\rho}$'s lie outside the interval $\pm 2/\sqrt{n}$ (a 95% confidence interval)
- One should expect approximately 1 out of 20 to lie outside the interval if the sequence is a white noise. Many more than that would invalidate the whiteness assumption.
- This allows for a recursive approach for manipulating / de-trending series until they are white noise, called **whitening**.
- The R function `acf` automatically displays those bounds with dashed blue lines.

SOI autocorrelation

```
acf(soi, main = "Sample autocorrelation function (SACF) of SOI")
```





```
r <- round(acf(soi, 6, plot = FALSE)$acf[-1], 3) # first 6 sample acf values
```

```
## [1] 0.604 0.374 0.214 0.050 -0.107 -0.187
```

The SOI series is clearly not a white noise.

```
plot(stats::lag(soi, -1), soi, main = "SOI pairs of values 1 month apart")  
legend("topleft", legend = r[1])  
plot(stats::lag(soi, -6), soi, main = "SOI pairs of values 6 months apart")  
legend("topleft", legend = r[6])
```



Scatterplots allow to have a visual representation of the dependence (which may not necessarily be linear).

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- Sample autocovariance function
- Sample autocorrelation function
- Testing for significance of autocorrelation
- **Sample cross-covariances and cross-correlations**
- Testing for independent cross-whiteness

Sample cross-covariances and cross-correlations

The **sample cross-covariance function** is

$$\hat{\gamma}_{xy}(h) = \frac{1}{n} \sum_{t=1}^{n-h} (x_{t+h} - \bar{x})(y_t - \bar{y}),$$

where $\hat{\gamma}_{xy}(-h) = \hat{\gamma}_{yx}(h)$ determines the function for *negative* lags.

The **sample cross-correlation function** is

$$-1 \leq \hat{\rho}_{xy}(h) = \frac{\hat{\gamma}_{xy}(h)}{\sqrt{\hat{\gamma}_x(0)\hat{\gamma}_y(0)}} \leq 1.$$

Note:

- Graphical examinations of $\hat{\rho}_{xy}(h)$ provide information about the leading or lagging relations in the data.

6 Estimation of correlation (TS 1.5)

- Background
- Sample mean
- Sample autocovariance function
- Sample autocorrelation function
- Testing for significance of autocorrelation
- Sample cross-covariances and cross-correlations
- Testing for independent cross-whiteness

Testing for independent cross-whiteness

If x_t and y_t are *independent* linear processes then the large sample distribution of $\hat{\rho}_{xy}(h)$ has mean 0 and

$$\sigma_{\hat{\rho}_{xy}} = \frac{1}{\sqrt{n}}$$


if at least one of the processes is independent white noise.

This is very useful, and adds to the toolbox:

- This provides feedback about the quality of our explanation of the relationship between both time series: if we have explained the trends and relationships between both processes, then their residuals should be independent white noise.
- After each improvement of our model, significance of the $\hat{\rho}_{xy}$'s of the residuals can be tested: if we have independent cross-whiteness then we have a good model. If the $\hat{\rho}_{xy}$'s are still significant (outside the boundaries) then we still have things to explain (to add).

SOI and recruitment correlation analysis

```
acf(soi, 48, main = "Southern Oscillation Index")  
acf(rec, 48, main = "Recruitment")  
ccf(soi, rec, 48, main = "SOI vs Recruitment", ylab = "CCF")
```



The SCCF (bottom) has a different cycle, and peak at $h = -6$
suggests SOI leads Recruitment by 6 months (negatively).

Idea of prewhitening

- to use the test of cross-whiteness one needs to “prewhiten” at least one of the series
- for the SOI vs recruitment example, there is strong seasonality which, upon removal, may whiten the series
- we look at an example here that looks like the SOI vs recruitment example, and show how this seasonality could be removed with the help of sin and cos functions

Example:



- Let us generate two series x_t and y_t , for $t = 1, \dots, 120$, independently as

$$x_t = 2 \cos\left(2\pi t \frac{1}{12}\right) + w_{t1} \quad \text{and} \quad y_t = 2 \cos\left(2\pi[t + 5] \frac{1}{12}\right) + w_{t2},$$

where $\{w_{t1}, w_{t2}; t = 1, \dots, 120\}$ are all independent standard normals.




- this generates the data and plots it:

```
set.seed(1492)
num <- 120
t <- 1:num
X <- ts(2 * cos(2 * pi * t/12) + rnorm(num), freq = 12)
Y <- ts(2 * cos(2 * pi * (t + 5)/12) + rnorm(num), freq = 12)
par(mfrow = c(1, 2), mgp = c(1.6, 0.6, 0), mar = c(3, 3, 1, 1))
plot(X)
plot(Y)
```



- looking at the ACFs one can see seasonality

```
par(mfrow = c(3, 2), mgp = c(1.6, 0.6, 0), mar = c(3, 3, 1, 1))  
acf(X, 48, ylab = "ACF(X)")  
acf(Y, 48, ylab = "ACF(Y)")  
ccf(X, Y, 24, ylab = "CCF(X,Y)")
```



- furthermore the CCF suggests cross-correlation even though the series are independent


- what we do now is to “prewhiten” y_t by removing the signal from the data by running a regression of y_t on $\cos(2\pi t)$ and $\sin(2\pi t)$ and then putting

$$\tilde{y} = y_t - \hat{y}_t,$$

where \hat{y}_t are the predicted values from the regression.

- in the R code below, Yw is \tilde{y}

```
par(mgp = c(1.6, 0.6, 0), mar = c(3, 3, 1, 1))
Yw <- resid(lm(Y ~ cos(2 * pi * t/12) + sin(2 * pi * t/12), na.action = NULL))
ccf(X, Yw, 24, ylab = "CCF(X,Yw)", ylim = c(-0.3, 0.3))
```



References I

Shumway, Robert H., and David S. Stoffer. 2017. *Time Series Analysis and Its Applications: With r Examples*. Springer.